

Session recap — Munder Difflin orchestration of US-29, Sun 04 – Mon 05 Oct 2026

15:09–03:52 AEST · about 12h 40m · feature branch for US-29 (local only, not pushed)

In one line: The Orchestrator ran the seven-stage US-29 workflow (interactive dashboard filtering and drill-down) across three agents to completion. Four live-test fix rounds and three independent code reviews ended with zero open findings. Nothing is pushed yet.

Timeline

Time	Phase	What happened
15:38	Restart and design	The first run was cancelled and restarted. Claude Coder ran kickoff, the human answered 14 questions, and the design document (SDD) was written by 16:20.
16:27	SDD review	The broken SDD reviewer was recreated as Pam. After a usage-limit pause, it returned "request changes" (4 required, 1 minor) at 20:14.
20:18	Build and code review	Claude Coder amended the SDD and built the feature (20:38). Code review asked for changes (20:45). Fixes were finished by 21:30.
21:31	Closure check and first recap	The reviewer was paused by its usage limit, and the human chose to wait. It approved at 01:15 with live checks still pending, and a first recap followed.
01:33	Live testing rounds	The human published and tested the dashboard. Filters and chart clicks failed, and were fixed over four rounds (02:28–03:34).
03:35	Final review and recap	Review round 2 found 2 issues, which were fixed by 03:42. Round 3 closed everything (03:45), and the final recap was written (03:52).

Human overviewer — direction and decisions

Direction given

- 15:38–16:01 — Follow the US-29 instructions document. Cancel the first run, restart Stage 1, and never let the Orchestrator answer human questions.
- 16:27–16:30 — Don't change the workflow; recover and reassign the SDD reviewer only.
- 21:34–21:37 — Asked whether the closure review could be skipped, then chose to **wait** for it.
- 01:33 — Keep the format-test dashboard, and asked for live-check instructions.
- 02:24 — Supplied a findings document from live testing.

Decisions made (all answered in Claude Coder's session)

Time	Decision	Why it mattered
16:15–16:17	Answered 14 kickoff questions, plus how to treat a teammate's change	Set the design scope
20:25–20:31	Answered Q16–Q20	Settled the SDD review fixes (Q16–Q18), how to get the dashboard format (Q19) and the table row count (Q20)
20:50–21:28	Q21 and Q22: tried the workspace "ignore filter" option, found it missing, and chose the fallback	Closed the last code-review finding

Time	Decision	Why it mattered
02:31	Q23 a, Q24 a, plus a new rule: remove any filter that errors or shows only blanks	Two empty filters were removed
02:40	Q25 a: bring the enrolment fields into the prediction data	Fixed the filter errors
03:20	Q26 a (AI's recommended option): point the charts at one data source	Fixed the chart-click errors
03:34	Q27 a (AI's recommended option): accept a row click narrowing to one student	Closed the last test row

Checks and corrections

- 15:58 — Cancelled the first run, because kickoff questions had been answered without the human.
- 16:27 — Stopped the Orchestrator from changing the workflow.
- 03:32 — Corrected a test report: the "B5" note was really about B7.

AI agents — workload

Orchestrator (main agent)

- Sent 22 work orders, from 15:39 to 03:45, one per stage or fix round, each with goal, output, tools and limits.
- Passed every human question across word for word, and checked each answer in Claude Coder's session before acting.
- Restarted two reviewers after usage-limit pauses (20:10 and 01:12). Recovered a reviewer report that had been lost to a file-format fault.
- Kept the task board, ASK ME questions and status board current, and ran the hourly standups.

Agents

Agent	Job	Ran	Result
Claude Coder	Design, build, tests, fixes, recaps	15:39–03:52	13 commits; 353 tests pass, 14 skipped; 15 of 15 deliberate faults caught; 2 recaps
Codex SDD Reviewer (Pam)	Independent design review	16:30–20:14	Request changes: 4 required, 1 minor; all fixed
Codex Code Reviewer	Code review plus 3 closure checks	20:39–03:45	Rounds: changes, approved pending live checks, changes, then approve with notes, 0 open

AI decisions

- After review round 2, the Orchestrator sent findings back to Claude Coder without asking, as the workflow requires.
- Claude Coder logged two agreed exceptions for changing a teammate's dataset and charts.

Outcomes

-  Design document written and independently reviewed; all findings fixed.
-  Dashboard filtering and drill-down built. The reviewer independently ran 353 passing tests.
-  All live test rows passed or were accepted by the human (one doesn't apply).
-  Final independent review: approved, zero open findings.
-  Live dashboard behaviour rests on the human's recorded results. The reviewer did not repeat them.

What's next

#	Task	Who	Status	Source
1	Approve the push and pull request to main	You	✅ Completed — you approved it; Claude Coder pushed the branch and opened pull request 132 (not merged)	Your instruction after this recap
2	Decide whether to commit the recaps	You	✅ Completed — you will add the recaps to the project files yourself; no agent commits them	Your instruction after this recap
3	Review and merge pull request 132	You	Open — waiting on your manual approval on GitHub	Claude Coder's pull-request report

Waiting on others: None.

You gave 12 prompts in the Orchestrator session and 28 question answers in Claude Coder's session. The AI agents made 13 commits, 4 review reports and 2 recaps.

Sources: sessions 0d54a2f4 (Orchestrator) and aeb5331d (Claude Coder); commits d2be96d-0877462; pull request karen-lee1029/Capstone_T115#132; review reports in the two reviewers' handoffs/US-29 folders; the Orchestrator's outbox log.